

Hadoop

Q) List the uses and responsibilities of following tools of Hadoop Ecosystem.

i) HDFS : (Hadoop distributed file system) is the primary storage system used by Hadoop to store large datasets across multiple machines in a distributed fashion.

Responsibilities:

1) Storage of large datasets: HDFS splits large datasets into blocks and stores them across different nodes for parallel processing.

2) Fault tolerance: HDFS replicates data across multiple nodes to ensure data is not lost in case of hardware failure.

3) Data Accessibility: It allows random access to stored data.

4) High throughput: High throughput than low latency access to large datasets, making it ideal for batch processing.

5) Scalability: HDFS handles datasets that grows exponentially.

ii) YARN (Yet Another Resource Negotiator) is Hadoop's resource management system that handles job scheduling and resource allocation across the Hadoop cluster. It allows multiple applications to run simultaneously by dynamically managing resources.

Responsibilities:

1) Resource Management: allocates resources like CPU and memory to various applications running on the Hadoop cluster.

iii) Job scheduling: It schedules jobs efficiently, managing the execution of tasks based on available resources.

iv) Multi tenancy: enables running different types of applications on the same cluster without conflict.

v) Fault tolerance: YARN tracks & restarts failed jobs, ensuring smooth execution of applications.

vi) Dynamic resource allocation: It ensures that resources are assigned to applications based on their current demands, improving cluster utilization.

vi) Flume: is a distributed, reliable and available service for efficiently collecting, aggregating, and moving large amount of log data or streaming data from various sources into centralized storage like HDFS.

Responsibilities:

1) Data Ingestion: Flume helps in ingesting data from multiple sources into HDFS.

2) Streaming data: This optimized for continuous real-time streams, handling large volumes of event data.

iii) Fault tolerance: Flume provides guaranteed delivery of data, even in case of node failure.

iv) Scalability: can be scaled horizontally by adding more agents to handle multiple applications to meet different data ingestion needs.

v) Customizability: Flume supports various custom sources, sinks & channel types to meet different data ingestion needs.

vi) Sequence : is a sequence of records using data flow modules like MySQL, PostgreSQL, Redis, etc.

Imports data from databases like MySQL, Oracle & PostgreSQL.

into HDFS and exports processed data back to the databases.
Responsibilities:

- i) Data Import: Sqoop imports structured data from relational databases into Hadoop for further analysis.
- ii) Data Export: It exports processed data from HDFS back to the relational database for reporting or further use.
- iii) Efficient Data Transfer: Sqoop uses parallel data from HDFS back to the relational database processing to efficiently transfer large datasets to and from Hadoop, ensuring high performance.
- iv) Data Transformation: During import/export, it can also transform the data into a desired format.
- v) Schema Mapping: Automatically generates schema for importing tables ensuring smooth integration with relational databases.

c) Zookeeper: is a centralized service for managing configuration information, naming, providing distributed synchronization and group services within Hadoop cluster. It is essential for coordinating distributed applications.

Responsibilities:

- i) Coordination: coordinates various distributed services & applications in a cluster to ensure they operate synchronously.
- ii) Configuration Management: Zookeeper stores & manages configuration information and distributes it across all nodes in the cluster.
- iii) Leader Election: facilitates leader election in a distributed system to ensure one master node controls task execution while others are standby.

- iv) Distributed Locking: provides a locking mechanism to avoid conflicts in accessing shared resources in a distributed environment.
- v) Fault Tolerance: Zookeeper itself is highly available and ensures consistency across nodes even if one or more nodes fail.

f) Hive: is a data warehouse infrastructure built on top of Hadoop that provides data summarization, querying, and analysis through an SQL-like interface (HiveQL), it allows users familiar with SQL to query data stored in HDFS without writing complex MapReduce programs.
Responsibilities:

- i) SQL-like Querying: Hive allows users to write queries in HiveQL to retrieve and manipulate large datasets in HDFS.
- ii) Data warehousing: Hive is primarily used for data warehousing tasks like data summarization, ad-hoc querying and analysis.
- iii) Schema on Read: Hive applies a schema to data as it is read (rather than when it is written), making it flexible for analyzing unstructured or semi-structured data.
- iv) Integration with Hadoop: Hive translates HiveQL queries into mapreduce jobs, allowing data to be processed on Hadoop's distributed computing framework.
- v) Extensibility: Hive supports user-defined functions (UDFs) for custom data processing & analysis, enabling users to extend its capabilities.

Q2) Explain the responsibilities of Name Node and Data Nodes in HDFS. Differentiate normal file system with HDFS in tabular column.

A) The Name Node is the central part of HDFS and acts as the master server responsible for managing the metadata and file system namespace. It coordinates the storage & retrieval of the data blocks from the data nodes.

Responsibilities of Name Node

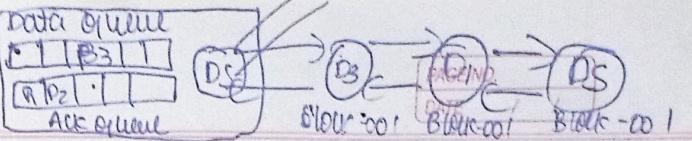
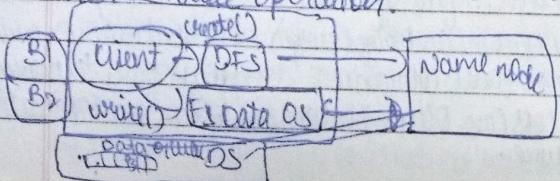
- Metadata Storage: Name Node stores metadata for the entire file system in memory.
- Namespace Management: It maintains the hierarchical structure of the file system, including directories, subdirectories, and files.
- Data Block Replication: It ensures data is fault tolerant by replicating blocks across multiple data nodes according to configured replication factor.
- Data Block Management: Name Node tracks the locations of data blocks stored on different data nodes, ensuring they are correctly replicated and available.
- Client Interaction: When a client wants to read or write data, the Name Node provides the necessary block info and coordinates access to data nodes.
- Failover Support: In case of data node failures, the Name Node ensures that the data is available by tracking replicas & initiating replication when necessary.

Data Nodes serve as the workers in HDFS, responsible for storing actual data blocks, they regularly communicate with the Name Node to ensure file system function. Responsibilities of Data Nodes:

- Storage of data blocks: Data nodes store and manage the physical storage of data blocks assigned by the Name Node.
- Data Block Reports: They periodically send block reports to the Name Node, indicating the blocks they are storing.
- Data Replication: Data Nodes participate in the replication process by replicating data to other Data Nodes as instructed by the Name Node.
- Data Read/Write: They manage the actual data I/O operations providing data to clients or accepting new data for storage.
- Heartbeat Signals: Data Nodes send heartbeat signals to the Name Node at regular intervals to indicate they are functioning properly.
- Error Handling: If a data node fails, its data blocks are replicated on other nodes by the Name node to ensure fault tolerance.

Difference Between Normal File System and HDFS	
Storage	Normal file system HDFS (Horizontally Distributed)
Storage	Single machine (centralized storage) Distributed across multiple.
Fault	Limited fault tolerance; High fault tolerance; data loss can occur on hardware failure.
	Data blocks are replicated.

Data size	Handles small to medium sized files.	Designed to handle very large datasets (GB, PB, PB+).
Scalability	Centralized file system, accessed by one machine.	Distributed file system, accessible across the cluster.
File size	Small block size (e.g. 4KB, 8KB)	Large block size (128MB or 256MB)
Type	Hierarchical file system	Designed for write Once, read many access pattern.
Redundancy	No automatic replication	Automatic duplication of data blocks for fault tolerance.
Scalability	Limited scalability as it's restricted to single machine storage capacity.	High scalability, can easily scale out by adding more data nodes.
Data flow	Data processing is done on the same machine where the file system resides.	Data is processed in a distributed manner across multiple nodes.
Handling	Stores metadata with each file locally.	Name node stores metadata in memory while data nodes store only data.



- i) Client request: sends a request to Name Node to create a file & write data, the Name Node checks if the file already exists and whether the client has the appropriate permissions to create the file.
 - ii) Block Assignment: The Name Node splits the file into blocks (128 MB) & identifies a list of Data Nodes to store each block. It returns the list Data Nodes where the replicas of the data blocks should be stored.
 - iii) Pipeline Setup: The Client establishes a pipeline connection with the first Data Node (Replica1). This Data Node, in turn, communicates with the second Data Node (Replica2), and the second Data Node communicates with the third Data Node. This is done in a pipeline fashion, ensuring all replicas are written simultaneously.
 - iv) Acknowledgment: After the block is written to all replicas, the Data Node acknowledges the successful write.
 - v) Writing Data: The client starts writing data blocks. Each block is written to Data Node.
 - vi) File Close: After all blocks are written & acknowledged, the client informs the Name Node that the file has been closed. The Name Node then commits the file creation.

Steps to recover from failure during HDFS write operation
Data node failure if node fails
The client is notified about the failure.

The Name node handles the failed Data node & managing the write. (Q4)

- * re-pipeline the write.

Client failure.

- * Name node waits for the client to re-establish the connection and continue the write operation.
- * If the client does not reconnect, the Name node marks the file as incomplete & removes written blocks.

Network failure.

- * The current block being written is saved in the existing Data nodes that have received the data.
- * Once the network recovers, the Name node re-establishes the connection with new Data nodes.

Name node failure:

- * HDFS uses a secondary Name node or Checkpoint node to recover metadata and file system state.
- * The system brings the Name node back online, and the incomplete write operations can be restarted from where they left off.

Handling data corruption uses mechanisms to verify data integrity.

Name node detects this using block reports. The corrupt block is deleted and a new replica is created from a healthy copy on another Data node.

Explain how HDFS ensures high availability of data & services. (Q4)

HDFS ensures high availability of data & services through a combination of redundancy, fault tolerance, and failover mechanisms.

- i) data Replication: block is replicated 3 times.
- v) primary replica, secondary replica tertiary replica
- ii) fault tolerance: if node fails, the Name node automatically redirects the client to another node.

Self healing mechanism.
iii) Name node High Availability. Traditionally HDFS had a single Name node, creating a single point of failure. However modern Hadoop version implements Name node High Availability (HA).

- iv) Active and Standby name node, failover mechanism.
- v) Fault Awareness to ensure replicas are placed on different rack.

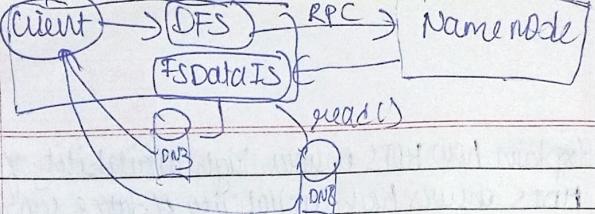
v) heartbeat & Block Report.

vi) checksum verification.

vii) Backup & Failover Mechanism.

(Q5) With a diagram show how read operation is performed in HDFS. Explain the steps involved in read operation.

- A) The read operation in HDFS involves interactions b/w the Client, Name node, Data nodes.

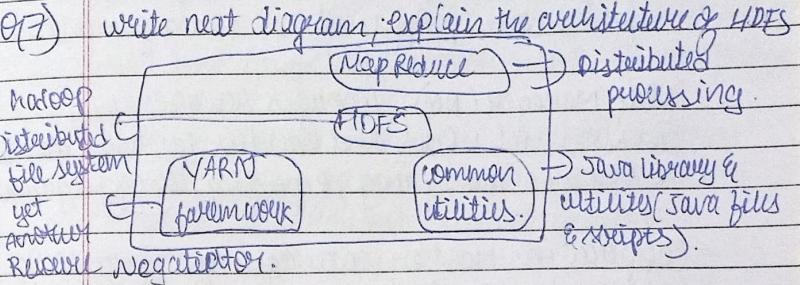


The check point process consolidates the in-memory metadata and the edit log into a single file system image.

PAGE NO. / /
DATE / /

media Recovery, efficiency,

- i) Client Request to Name node
- ii) block location: Retrieved from Name node.
- iii) connecting to data nodes.
- iv) Data Block Reading.
- v) Block replicates as Backup.
- vi) Data return to client
- vii) No interaction with name node during data transfer.



Client interacts with HDFS to perform read & write.
HDFS → Name node (master node)
Data node (slave node).

blocks (data unit of storage in HDFS (128 MB))
Map reduce → blocks of data are mapped and shuffle & sort then reduced by reducers.
YARN → Resource Negotiator & job scheduling Resource management.

- Q6) Explain the terms heart beat, edit log and check points.
 - 1) heartbeat: periodic signal sent by a Data node to the Name node to indicate that it is alive & functioning correctly.
 - 2) Health monitoring, Data node availability mechanism: Data nodes send heart beats at regular intervals (default is every 3 seconds).
 - 3) edit log: is a file maintained by the Name node that records all changes made to the file system metadata. It logs operations such as file creation, deletion and modification.
 - 4) Check point: is a snapshot of the file system metadata and its state, including the edits recorded in the edit log.

- Q7) Write next diagram, explain the architecture of HDFS.
 - MapReduce
 - distributed processing
 - HDFS
 - YARN
 - common utilities (Java files & scripts)
 - Yet Another Resource Negotiator
 - Java library & utilities (Java files & scripts)
- Client interacts with HDFS to perform read & write.
HDFS → Name node (master node)
Data node (slave node).
- blocks (data unit of storage in HDFS (128 MB))
Map reduce → blocks of data are mapped and shuffle & sort then reduced by reducers.
YARN → Resource Negotiator & job scheduling Resource management.
- Q8) List the major components of YARN. Explain the responsibilities of YARN.
- YARN (Yet Another Resource Negotiator) is a resource management layer of Hadoop that allows for more flexible & efficient resource management compared to the other map reduce frameworks.

Resource Manager allocates cluster resources.

- ↳ keeps track of available resources
- ↳ scheduler, Application management.

Node Manager (NM): manages resources on individual nodes and handles container creation.

- ↳ Resource monitoring, Reporting, log aggregation

Application Master: manages the lifecycle of a specific application & coordinate with the resource manager & nodemanager

Responsibilities of YARN

- ↳ Resource allocation, Application scheduling
- ↳ Resource monitoring, fault tolerance & recovery
- ↳ Cluster utilization, Scalability

(Q9) list out the responsibilities of Resource Manager

- ↳ Resource allocation & Management.
- ↳ Application submission & Management
- ↳ Application master coordination
- ↳ Resource monitoring, Reporting
- ↳ failure detection & recovery
- ↳ queue Management & Scalability
- ↳ load balancing & integration with other components.

(Q10)

Explain different schedulers used in YARN

- i) capacity Scheduler ensures that each queue gets its fair share of resources and that large applications do not monopolize cluster.
- ii) fair scheduler aims to allocate resources equally among running applications to ensure that all applications receive fair share
- iii) FIFO is the simplest form of scheduling in YARN

(Q11)

Explain different types of failure in mapreduce job.

- i) Task failure occurs when an individual map or reduce task fails to complete successfully.
- ii) Node failure occurs when a node in the hadoop cluster becomes unresponsive or crashes.
- iii) Resource Manager failure: Can disrupt the management of job scheduling & resource allocation.
- iv) Data Node failure occurs when a datanode failing the data blocks become unresponsive.
- v) Application Master failure occurs when the application master, which manages the lifecycle of jobs.
- vi) Network failure can cause communication issues between nodes, leading to task failure by delay.

- (12) Give different reasons for task failure in map reduce job give the steps for recovering from Application master failure.
- Hardware Failure * Disk Failure * Network Issues
 - Software Bugs: Application Bugs * Library or Dependency Issues
 - Data Corruption: Corrupted Input Data * Data Format Issues
 - Resource Constraints - CPU or memory limitations, insufficient Disk Space
 - Network Issues: network partitions, slow network
 - Configuration Errors: misconfiguration, resource misallocation
 - Dependency failure: external service failure
 - File System Issues: HDFS failure

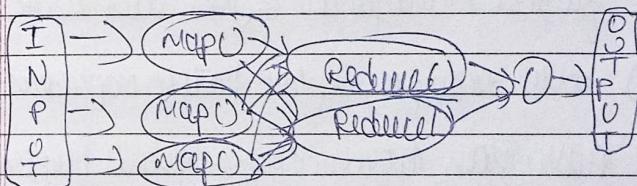
Steps for Recovering from Application Master

- Detect failure
- Restart Application Master
- Recover Application State
- Handle Data and Resource Allocation
- Notify & Update Users
- Post failure Actions

(Q13)

With neat diagram, show how map reduce jobs are executed in Hadoop environment.

- * Job Submission: user submits job.
- * Job Initialization: creates a job context
- * Job Scheduling: schedules the job's tasks (map and reduce) across the cluster nodes
- * Map Task Execution: input split, applies map()
- * Shuffling & sort: transfer the started map output to the appropriate reduce tasks.
- * Reduce Task Execution: processes intermediate key-value pairs, applies the reduce function
- * Job Completion: job tracker to resource manager * Job status.



Give the responsibilities of

- Application Master
- Application Manager
- Name Manager
- Resource Manager
- Application Master: Application life cycle management, Resource requests & Negotiation, Task Scheduling, Application Monitoring

Application Master or manages the Application Master and acts as a middle ware for the Resource Manager & Application Master. If any Application Master is failed then it restarts the execution, re compete that.

○ Node Manager : Resource monitoring, Container Management, Task execution, Health reporting
\\$ Log Aggregation

i) Resource Manager : oversees resource allocation, scheduler management, resource tracking, application coordination, failure management & job submission handling

Ques) with example, explain Hadoop MapReduce process

Public class WordCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {

String line = value.toString();

StringTokenizer tokenizer = new StringTokenizer(line);

while (tokenizer.hasMoreTokens()) {

word.set(tokenizer.nextToken());

context.write(word, one);

public class WordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
int sum = 0;
for (IntWritable val : values) {
sum += val.get();
}
context.write(key, sum);
}

Count the number of occurrences of each word in a large collection of text files mapped by process.

▷ Job Submission, Job Initialization.
Map Phase - shuffle and sort phase,
Reduce phase - on completion.