


Logistic Regression


Students' Acceptance at a University



Students' Acceptance at a University



Logistic Regression



Logistic Regression

Machine Learning Approach

Logistic regression

- Logistic regression is one of the most popular machine learning algorithms for **binary classification**.
- It is a simple algorithm that performs very well on a wide range of problems.

Linear Regression:


- **Simple:**



$$y = b_0 + b_1 * x$$


- **Multiple:**


$$y = b_0 + b_1 * x_1 + \dots + b_n * x_n$$

We know this:











$$y = b_0 + b_1 * x$$


Sigmoid Function

$$p = \frac{1}{1 + e^{-y}}$$


$$\ln \left(\frac{p}{1 - p} \right) = b_0 + b_1 * x$$




y (Actual DV)




y (Actual DV)







\hat{p} (Probability)



\hat{y} (Predicted DV)



\hat{y} (Predicted DV)




Dataset

- Dataset has two input variables (X1 and X2) and one output variable (Y).
- The input variables are real-valued random numbers drawn from a Gaussian distribution.
- The output variable has two values, making the problem a binary classification problem.

Input	Input	Actual
X1	X2	Output Y
2.7810836	2.550537003	0
1.465489372	2.362125076	0
3.396561688	4.400293529	0
1.38807019	1.850220317	0
3.06407232	3.005305973	0
7.627531214	2.759262235	1
5.332441248	2.088626775	1
6.922596716	1.77106367	1
8.675418651	-0.2420686549	1
7.673756466	3.508563011	1

Plot of the Dataset

- we can easily draw a line to separate the classes. we are going to do with the logistic regression model.



Input	Input	Actual
X1	X2	Output Y
2.7810836	2.550537003	0
1.465489372	2.362125076	0
3.396561688	4.400293529	0
1.38807019	1.850220317	0
3.06407232	3.005305973	0
7.627531214	2.759262235	1
5.332441248	2.088626775	1
6.922596716	1.77106367	1
8.675418651	-0.2420686549	1
7.673756466	3.508563011	1

Logistic Function

- The logistic function is the heart of the logistic regression technique.
- The logistic function is defined as:

$$\text{transformed} = 1 / (1 + e^{-x})$$


Where **e** is the numerical constant Euler's number and **x** is a input we plug into the function.

Logistic Function

- Let's plug in a series of numbers from -5 to +5 and see how the logistic function transforms them.
- You can see that all of the inputs have been transformed into the range [0, 1]
- The smallest negative numbers resulted in values close to zero
- The larger positive numbers resulted in values close to one.
- See that 0 transformed to 0.5 or the midpoint of the new range.

X	Transformed
-5	0.006692850924
-4	0.01798620996
-3	0.04742587318
-2	0.119202922
-1	0.2689414214
0	0.5
1	0.7310585786
2	0.880797078
3	0.9525741268
4	0.98201379
5	0.9933071491

$$\text{transformed} = 1 / (1 + e^{-x})$$



Logistic Regression Model

- The logistic regression model takes real-valued inputs and makes a prediction


Prediction

If **transformed** ≥ 0.5

$y_{\text{pred}} = \text{class 1}$

otherwise

$y_{\text{pred}} = \text{class 0}$



Logistic Regression Model

- For this dataset, the logistic regression has three coefficients:

$$\mathbf{x} = \mathbf{b0} + \mathbf{b1} * \mathbf{x_1} + \mathbf{b2} * \mathbf{x_2}$$

- The job of the learning algorithm will be to discover the best values for the coefficients (b0, b1 and b2) based on the training data.
- The output is transformed into a probability using the logistic function:

$$\text{transformed} = 1 / (1 + e^{(-x)})$$

Input X_1	Input X_2	A_Output Y
2.7810836	2.550537003	0
1.465489372	2.362125076	0
3.396561688	4.400293529	0
1.38807019	1.850220317	0
3.06407232	3.005305973	0
7.627531214	2.759262235	1
5.332441248	2.088626775	1
6.922596716	1.77106367	1
8.675418651	-0.2420686549	1
7.673756466	3.508563011	1

Calculate Prediction

Let's start off by assigning 0.0 to each coefficient and calculating the probability of the first training instance that belongs to class 0.

$b_0 = 0.0; b_1 = 0.0; b_2 = 0.0$

The first training instance is: $x_1=2.7810836$, $x_2=2.550537003$, $x=0$

$\text{transformed} = 1 / (1 + e^{-x})$

Using the above equation - calculate a prediction

$$\begin{aligned}\text{transformed} &= 1 / (1 + e^{-(b_0 + b_1 * x_1 + b_2 * x_2)}) \\ &= 1 / (1 + e^{-(0.0 + 0.0 * 2.7810836 + 0.0 * 2.550537003)})\end{aligned}$$

$$\text{transformed} = 0.5$$

Calculate the new coefficient values using a simple update equation

$$b(\text{new}) = b(\text{old}) + \text{alpha} * (x - \text{transformed}) * \text{transformed} * (1 - \text{transformed}) * x_i$$

Let's update the coefficients using the prediction (0.5) and coefficient values (0.0) from the previous section.

$$b_0 = 0 + 0.3 * (0 - 0.5) * 0.5 * (1 - 0.5) * 1.0 = -0.0375$$

$$b_1 = 0 + 0.3 * (0 - 0.5) * 0.5 * (1 - 0.5) * 2.7810836 = -0.104290635$$

$$b_2 = 0 + 0.3 * (0 - 0.5) * 0.5 * (1 - 0.5) * 2.550537003 = -0.09564513761$$

Calculate output using new b0, b1 and b2

$$x = b_0 + b_1 * x_1 + b_2 * x_2$$

$$= -0.0375 - 0.104290635 * 2.7810836 - 0.09564513761 * 2.550537003$$

$$= -0.5565$$

$$\text{transformed} = 1 / (1 + e^{-x})$$

$$= 1 / (1 + e^{-(-0.5565)})$$

$$= 0.364$$

If transformed ≥ 0.5

y_pred = class-1

otherwise

y_pred = class-0

Repeat the Process

- We can **repeat this process** and update the model **for each training instance** in the dataset.
- A single iteration through the training dataset is called an epoch. It is common to repeat the stochastic gradient descent procedure for a fixed number of epochs.
- At the end of epoch you can calculate error values for the model. Because this is a classification problem, it would be nice to get an idea of how accurate the model is at each iteration.
- The graph below show a plot of accuracy of the model **over 10 epochs**.

Input X1	Input X2	Actual Output Y
2.7810836	2.550537003	0
1.465489372	2.362125076	0
3.396561688	4.400293529	0
1.38807019	1.850220317	0
3.06407232	3.005305973	0
7.627531214	2.759262235	1
5.332441248	2.088626775	1
6.922596716	1.77106367	1
8.675418651	-0.2420686549	1
7.673756466	3.508563011	1

transformed =
 $1 / (1 + e^{-\text{output}})$

P_output =
 if p(class) < 0.5) then 0 else 1

0.2987569857	0
0.145951056	0
0.08533326531	0
0.2197373144	0
0.2470590002	0
0.9547021348	1
0.8620341908	1
0.9717729051	1
0.9992954521	1
0.905489323	1

output = b0 + b1*x1 + b2*x2


Compare actual & prediction

- You can see that the model very quickly achieves 100% accuracy on the training dataset.
- The coefficients calculated after 10 epochs of stochastic gradient descent are:

$b_0 = -0.4066054641$

$b_1 = 0.8525733164$

$b_2 = -1.104746259$



Calculate Accuracy

- Finally, we can calculate the accuracy for the model on the training dataset:

$$\text{accuracy} = (\text{correct predictions} / \text{num predictions made}) * 100$$

$$\text{accuracy} = (10 / 10) * 100$$

$$\text{accuracy} = 100\%$$

Thank you