


# Reinforcement Learning

BDA/ HDA

# What Is Machine Learning?

*Machine learning is a subset of artificial intelligence (AI) which provides machines the ability to learn automatically & improve from experience without being explicitly programmed.*



# Types Of Machine Learning



Supervised Learning



Unsupervised Learning




Reinforcement Learning

# What Is Reinforcement Learning?

*Reinforcement learning is a type of Machine Learning where an agent learns to behave in a environment by performing actions and seeing the results*



Agent




Environment



Reward


# Reinforcement Learning With An Analogy

---



# Reinforcement Learning With An Analogy

---



# Reinforcement Learning vs. Supervised Learning



Parameters	Reinforcement Learning	Supervised Learning
Decision style	Take <b>decisions sequentially</b> .	Decision is made on <b>the input given at the beginning</b> .
Works on	Works on <b>interacting with the environment</b> .	Works on examples or given <b>sample data</b> .
Dependency on decision	Decision is dependent.	Decisions are independent of each other.
Best suited	Supports and work better in AI, where human interaction is prevalent.	It is mostly operated with an interactive software system or applications.
Example	Chess game	Object recognition

# Reinforcement Learning Process


Reinforcement Learning system is comprised of two main components:

- Agent
- Environment

Agent



Environment



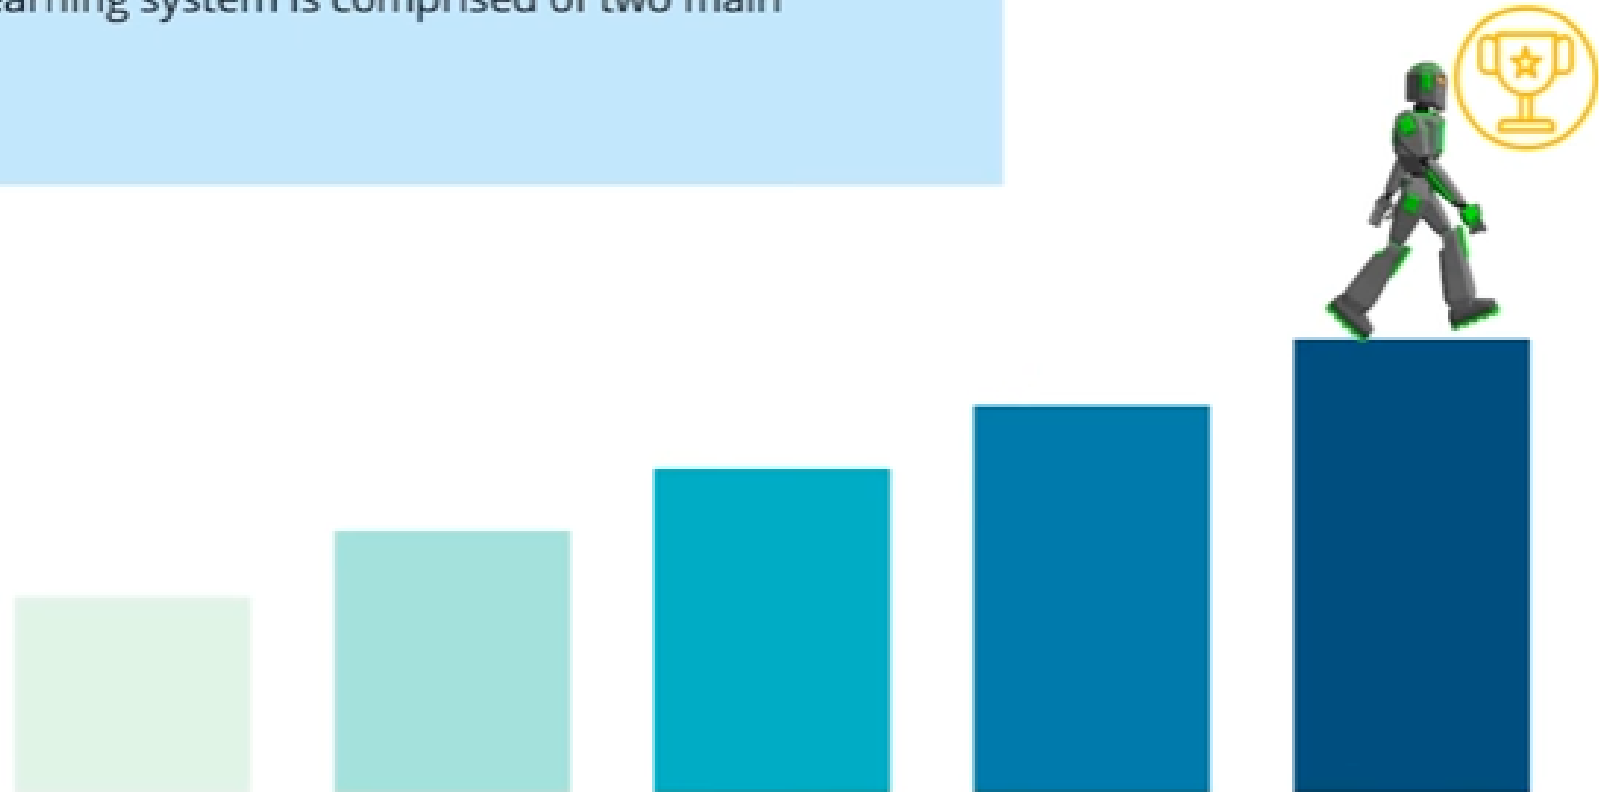


# Reinforcement Learning Process

---


Reinforcement Learning system is comprised of two main components:

- Agent
- Environment



# Counter Strike Example

---



1. The RL Agent (Player1) collects state  $S^0$  from the environment
2. Based on the state  $S^0$ , the RL agent takes an action  $A^0$ , initially the action is random
3. The environment is now in a new state  $S^1$
4. RL agent now gets a reward  $R^1$  from the environment
5. The RL loop goes on until the RL agent is dead or reaches the destination

# Reinforcement Learning Definitions

---



Agent: The RL algorithm that learns from trial and error

Environment: The world through which the agent moves



Action (A): All the possible steps that the agent can take

State (S): Current condition returned by the environment



# Reinforcement Learning Definitions



Reward ( $R$ ): An instant return from the environment to appraise the last action



Policy ( $\pi$ ): The approach that the agent uses to determine the next action based on the current state




Value ( $V$ ): The expected long-term return with discount, as opposed to the short-term reward  $R$



Action-value ( $Q$ ): This similar to Value, except, it takes an extra parameter, the current action ( $A$ )

# Reward Maximization


Reward maximization theory states that, *a RL agent must be trained in such a way that, he takes the best action so that the reward is maximum.*



# Exploration & Exploitation

*Exploitation* is about using the already known exploited information to heighten the rewards

*Exploration* is about exploring and capturing more information about an environment



Agent



Opponent




Reward

# Markov Decision Process

The mathematical approach for mapping a solution in reinforcement learning is called *Markov Decision Process* (MDP)

The following parameters are used to attain a solution:

- Set of actions,  $A$
- Set of states,  $S$
- Reward,  $R$
- Policy,  $\pi$
- Value,  $V$




# Markov Decision Process – Shortest Path Problem

Goal: Find the shortest path between A and D with minimum possible cost (maximum reward)

In this problem,


- Set of states are denoted by nodes i.e. {A, B, C, D}
- Action is to traverse from one node to another {A → B, C → D}
- Reward is the cost represented by each edge
- Policy is the path taken to reach the destination {A → C → D}





# Understanding Q-Learning With An Example

*Place an agent in any one of the rooms (0,1,2,3,4) and the goal is to reach outside the building (room 5)*




- 5 rooms in a building connected by doors
- each room is numbered 0 through 4
- The outside of the building can be thought of as one big room (5)
- Doors 1 and 4 lead into the building from room 5 (outside)

# Understanding Q-Learning With An Example

*Next step is to associate a reward value to each door:*


- doors that lead directly to the goal have a reward of 100
- Doors not directly connected to the target room have zero reward
- Because doors are two-way, two arrows are assigned to each room
- Each arrow contains an instant reward value



# Understanding Q-Learning With An Example

The terminology in Q-Learning includes the terms state and action:

- Room (including room 5) represents a state
- agent's movement from one room to another represents an action
- In the figure, a state is depicted as a node, while "action" is represented by the arrows




Example (Agent traverse from room 2 to room5):

1. Initial state = state 2
2. State 2  $\rightarrow$  state 3
3. State 3  $\rightarrow$  state (2, 1, 4)
4. State 4  $\rightarrow$  state 5

# Understanding Q-Learning With An Example

We can put the state diagram and the instant reward values into a reward table, matrix  $R$ .



State

Action

	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

$R =$

The -1's in the table represent null values

# Understanding Q-Learning With An Example

Add another matrix Q, representing the memory of what the agent has learned through experience.

- The rows of matrix Q represent the current state of the agent
- columns represent the possible actions leading to the next state
- Formula to calculate the Q matrix:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max} [Q(\text{next state}, \text{all actions})]$$

## Note

The Gamma parameter has a range of 0 to 1 ( $0 \leq \text{Gamma} < 1$ ).

- If Gamma is closer to zero, the agent will tend to consider only immediate rewards.
- If Gamma is closer to one, the agent will consider future rewards with greater weight



# Q – Learning Example

First step is to set the value of the learning parameter  $\text{Gamma} = 0.8$ , and the initial state as Room 1.

Next, initialize matrix  $Q$  as a zero matrix:


- From room 1 you can either go to room 3 or 5, let's select room 5.
- From room 5, calculate maximum  $Q$  value for this next state based on all possible actions:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$

$$Q(1,5) = R(1,5) + 0.8 * \text{Max}[Q(5,1), Q(5,4), Q(5,5)] = 100 + 0.8 * 0 = 100$$

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$



# Q – Learning Algorithm

- 1 Set the gamma parameter, and environment rewards in matrix R
- 2 Initialize matrix Q to zero
- 3 Select a random initial state
- 4 Set initial state = current state
- 5 Select one among all possible actions for the current state
- 6 Using this possible action, consider going to the next state
- 7 Get maximum Q value for this next state based on all possible actions
- 8 Compute:  $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$
- 9 Repeat above steps until current state = goal state

# Q – Learning Example

First step is to set the value of the learning parameter  $\text{Gamma} = 0.8$ , and the initial state as Room 1.

Next, initialize matrix  $Q$  as a zero matrix:

- From room 1 you can either go to room 3 or 5, let's select room 5.
- From room 5, calculate maximum  $Q$  value for this next state based on all possible actions:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$


$$Q(1,5) = R(1,5) + 0.8 * \text{Max}[Q(5,1), Q(5,4), Q(5,5)] = 100 + 0.8 * 0 = 100$$

$Q =$

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	100
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

$R =$

State	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100





# Q – Learning Example

For the next episode, we start with a randomly chosen initial state, i.e. state 3


- From room 3 you can either go to room 1,2 or 4, let's select room 1.
- From room 1, calculate maximum Q value for this next state based on all possible actions:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$

$$Q(3,1) = R(3,1) + 0.8 * \text{Max}[Q(1,3), Q(1,5)] = 0 + 0.8 * [0, 100] = 80$$

The matrix Q get's updated

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$
$$R = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$



Thank you